

M23CS7.501:  
Advanced Natural Language Processing

Mid Semester Exam

Sep 22, 2023

MM: 30

Time: 1.5 hrs

Note: Marks are mentioned next to the questions.

1. Explain how RNNs enable handling long distance dependencies. Can you compare LSTMs and GRUs at an architectural level?  
[5+5 marks]
2. Describe the ELMO Contextual embedding model. How is it used for downstream tasks?  
[5+5 marks]
3. Describe a sequence-to-sequence machine translation model where decoder pays attention to the encoder. Why is the attention needed?  
[3 + 7 marks]